

Scientific computing using Cython: Best of both worlds!

November 5, 2017
Pycon India 2017

About me

Simmi Mourya

Deep learning Engineer, Predible
Health

<http://prediblehealth.com/>

Github: @simmimourya1

twitter: @simmimourya

What this talk is about?

What is Cython?

Unleash Cython superpowers.

Build your first cython module.

Cython and NumPy

Unleash more superpowers.

Project demo: Cyvlfeat

What is Cython?

Optimising static compiler for both the Python and Cython

It makes writing C extensions for Python as easy as Python itself.

UNLEASH SUPERPOWERS.



SUPERPOWERS

I HAS THEM

-
- Write Python code that calls back and forth from and to C or C++ code natively at any point.

-
- Easily tune readable Python code into plain C performance by adding static type declarations.

Pure Python code:

```
def f(x):  
    return x**2-x
```

```
def integrate_f(a, b, N):  
    s = 0  
    dx = (b-a)/N  
    for i in range(N):  
        s += f(a+i*dx)  
    return s * dx
```


With additional
Static Type
declarations: 4
times speedup

```
def f(double x):  
    return x**2-x
```

```
def integrate_f(double a, double  
b, int N):  
    cdef int i  
    cdef double s, dx  
    s = 0  
    dx = (b-a)/N  
    for i in range(N):  
        s += f(a+i*dx)  
    return s * dx
```

Def

def - Basically, it's Python

def is used for code that will be:

- Called directly from Python code with Python objects as arguments.
- Returns a Python object

Cdef

cdef - Basically, it's C

cdef is used:

- Where Cython functions are intended to be pure 'C' functions.
- All types *must* be declared.
- The generated code is about as fast as you can get though.

Cpdef

cpdef - It's both

cpdef combines both `def` and `cdef` by creating two functions; a `cdef` for C types and a `def` for Python types.

- Uses early binding when using C fundamental types
- Uses dynamic binding when Python objects are passed

→ It can call C/C++ libraries directly using
cimport

```
In [11]: %%cython -a
          # libc math functions
          from libc cimport math

          print( math.sin(math.M_PI / 2) )

1.0
```

```
Out[11]: Generated by Cython 0.23.beta1

Yellow lines hint at Python interaction.
Click on a line that starts with a "+" to see the C code that Cython generated for it.
1: # libc math functions
2: from libc cimport math
3:
+4: print( math.sin(math.M_PI / 2) )
```

→ Integrate natively with existing code.

A meme featuring a tabby cat with its front paws raised, emitting colorful laser beams (yellow, green, blue, and purple) that point towards a city skyline at night. The cat has a determined expression. The text "THE FORCE IS STRONG WITH THIS ONE." is overlaid in the center.

THE FORCE IS STRONG WITH THIS ONE.

Where's the magic?



Motivation

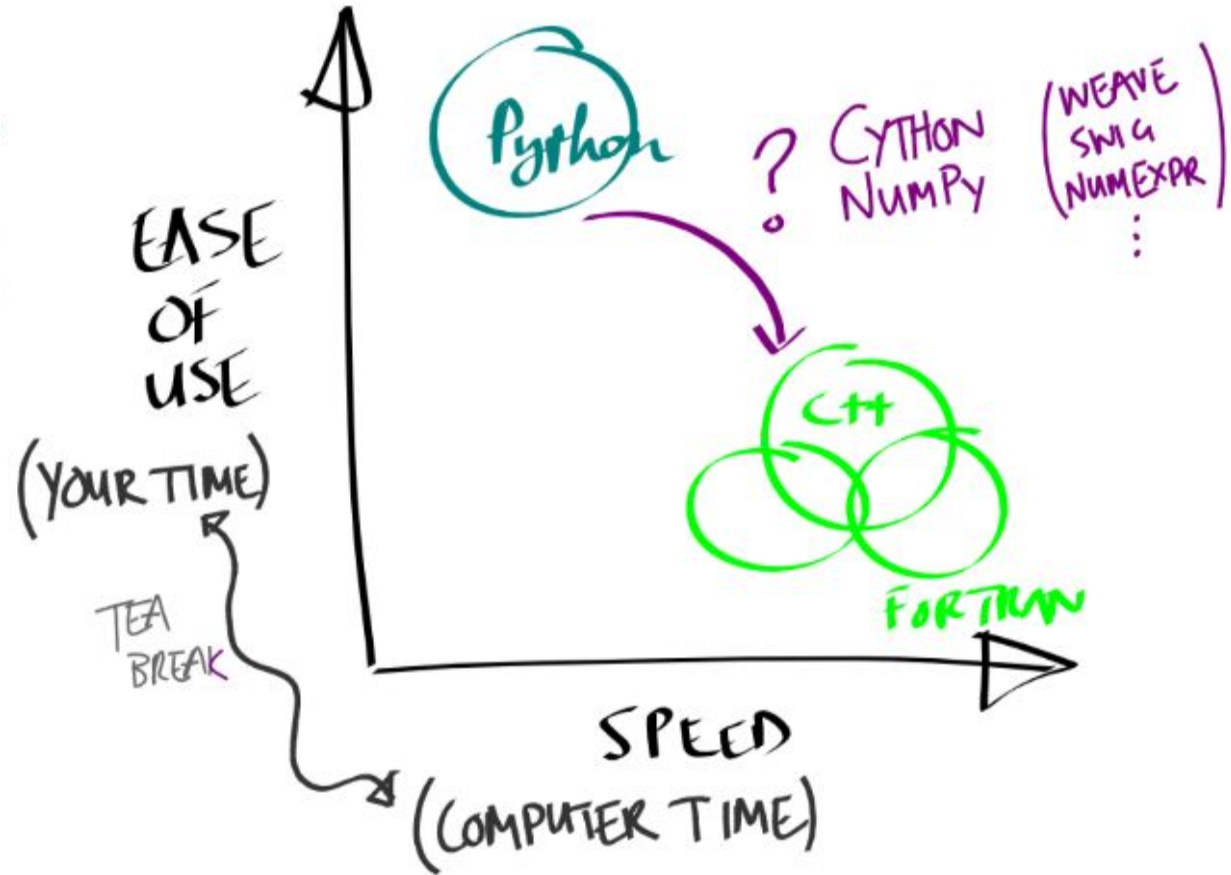


Illustration borrowed from: Stéfan van der Walt's presentation at Advanced Python Summer School, Kiel 2012.

Cython by example

Python

1x

```
def fib(n):  
    a,b = 1,1  
    for i in range(n):  
        a, b = a+b, a  
    return a
```

C/C++

100x

```
int fib(int n)  
{  
    int tmp, i, a, b;  
    a = b = 1;  
    for(i=0; i<n; i++){  
        tmp = a;  
        a += b;  
        b = tmp;  
    }  
}
```

Cython

80x

```
def fib(int n):  
    cdef int i, a, b  
    a, b = 1,1  
    for i in range(n):  
        a, b = a+b, a  
    return a
```

Cython in wild

Project	Cython files	Cython SLOC
sage	761	477,000
numpy	14	5,000
scipy	28	24,000
pandas	21	27,000
lxml	12	22,000
scikits-learn	35	15,000
scikits-image	48	11,000
mpi4py	48	12,000
yt	45	18,000

Projects master branches as of November 2014

A close-up photograph of a kitten's face. The kitten has light blue eyes and is looking directly at the camera. Its mouth is wide open, showing its pink tongue and teeth. The kitten's fur is a mix of white and brown. The background is dark and out of focus.

CODING IN CYTHON IS LIKE CODING IN
PYTHON AND C AT THE SAME TIME!

Use Case 1: Library Wrapping

Cython is a popular choice for writing Python interface modules for C libraries

Use Case 2: Performance-critical code

- Python
- High-level
- Slow
- No variables typed
- C/C++/Fortran
- Lower-level
- Fast
- All variables typed

Common procedure: Where speed is needed, use a compiled language, then wrap the code for use from Python.

Use Case 3: Breaking out of the GIL

Cython As a few of you might know, C Python has an infamous Global Interpreter Lock (GIL)

- It limits thread performance

The Unwritten Rules of Python

- You do not talk about the GIL.
- You do NOT talk about the GIL.

**Don't even
mention the
GIL.
No seriously...**



Consider this simple script, that runs twice, sequentially, a `busy_sleep`, i.e. a function that simulate a CPU intensive task:

```
from def busy_sleep(n):  
    while n > 0:  
        n -= 1  
N = 99999999  
busy_sleep(N)  
    busy_sleep(N)
```

This takes 6.7 seconds to run.

Now consider the threaded version:

```
from threading import Thread
def busy_sleep(n):
    while n > 0:
        n -= 1
N = 999999999
t1 = Thread(target=busy_sleep, args=(N, ))
t2 = Thread(target=busy_sleep, args=(N, ))
t1.start()
t2.start()
t1.join()
t2.join()
```

This takes 11.1 seconds to run. What?

Cython offers a wonderful context manager to run instructions without the GIL: *with nogil*.

Demo

With Nogil

```
def busy_sleep(int n):  
    _busy_sleep(n)  
  
def busy_sleep_nogil(int n):  
    with nogil:  
        _busy_sleep(n)  
  
cdef inline void _busy_sleep(int n) nogil:  
    cdef double tmp = 0.0  
    while n > 0:  
        tmp = (n ** 0.5) ** 0.5  
        n -= 1
```



Building Cython code

Building Cython code

Cython code must, unlike Python, be compiled.
This happens in two stages:

- A .pyx file is compiled by Cython to a .c file, containing the code of a Python extension module
- The .c file is compiled by a C compiler into a .so file (or .pyd on Windows) which can be imported directly into a Python session.

Building Cython code

There are several ways to build Cython code:

- Write a distutils setup.py.
- Use pyximport, importing Cython .pyx files as if they were .py files (using distutils to compile and build in the background).
- Use the Jupyter notebook or the Sage notebook, both of which allow Cython code inline.

Building Cython modules using distutils

Imagine a simple “hello world” script in a file `hello.pyx`:

```
def say_hello_to(name):  
    print("Hello %s!" % name)
```

The following could be a corresponding `setup.py` script:

```
from distutils.core import setup  
from Cython.Build import cythonize
```

```
setup(  
    name = 'Hello world app',  
    ext_modules = cythonize("hello.pyx"),  
)
```

Building Cython modules using distutils

To use this to build your Cython file use the command line options:

```
$ python setup.py build_ext --inplace
```

It leaves a file in your local directory called helloworld.so (unix) or helloworld.pyd (Windows).

```
>>> import helloworld  
Hello World
```

pyximport: Cython Compilation the Easy Way

To load .pyx files directly on import, without having to write a setup.py file.

```
>>> import pyximport; pyximport.install()  
>>> import hello  
Hello World
```

Build with Jupyter

Load the cython magic extension

```
In [1]: %load_ext cython
```

Then simply use the cython magic function to start writing cython code

```
In [2]: %%cython

cdef int a = 0
for i in range(10):
    a += i
print(a)
```

45

Build with Jupyter

Add --annotation or -a for showing the code analysis of the compiled code

```
In [3]: %%cython --annotate
```

```
cdef int a = 0
for i in range(10):
    a += i
print(a)
```

45

```
Out[3]:
```

Generated by Cython 0.25.2

Yellow lines hint at Python interaction.

Click on a line that starts with a "+" to see the C code that Cython generated for it.

```
1:
+2: cdef int a = 0
+3: for i in range(10):
+4:     a += i
+5: print(a)
```

C arrays

```
In [21]: %%cython -a
def carrays():
    cdef int[10] a, b
    a[:5] = [1,2,3,4,5]
    b = a
    b[5:] = [6,7,8,9,10]

    for i in b[:3]:
        print(i+1)

    return b
```

Out[21]: Generated by Cython 0.23.beta1

Yellow lines hint at Python interaction.

Click on a line that starts with a "+" to see the C code that Cython generated for it.

```
+01: def carrays():
    02:     cdef int[10] a, b
+03:     a[:5] = [1,2,3,4,5]
+04:     b = a
+05:     b[5:] = [6,7,8,9,10]
    06:
+07:     for i in b[:3]:
+08:         print(i+1)
    09:
+10:     return b
```

The conclusions:

1. Naive Cython does speed things up, but not by much (x1.8)
2. Optimised Cython is fairly effortless (in this case) and worthwhile (x2.5)
3. cpdef gives a good improvement over def
4. cdef is really valuable (x72)
5. Cython's cdef is insignificantly different from the more complicated C extension that is our best attempt

NumPy and Cython

Provides fast access to NumPy arrays

It has a C-level type -> *typed memoryview*

Similar to NumPy array buffer support.

Allows us to work with buffers without having to know the details.



Example: Typed Memoryviews

The `double[:] foo` syntax declares `foo` to be a typed memoryview.

```
def sum_arg(double[:] foo):  
    """Sums its argument's contents."""  
    cdef double i, total = 0.0  
    for i in foo:  
        total += i  
    return total
```

Providing C-Level access to Typed Memoryview data

```
def sum_arg(double[:] foo):  
    """Sums its argument's contents."""  
    cdef:  
        double total = 0.0  
        int j, length  
        length = foo.shape[0]  
    for j in range(length):  
        total += foo[j]  
    return total
```



Trading Safety for performance



Trading Safety for performance

```
from cython cimport boundscheck, wraparound
def sum_arg(double[:] foo):
    ...
    ...

    with boundscheck(False), wraparound(False)
        for j in range(length):
            ...
```

Trading Safety for performance

```
from cython cimport boundscheck, wraparound
@boundscheck(False)
@wraparound(False)

def sum_arg(double[:] foo):
    ...
    ...

    with boundscheck(False), wraparound(False)
        for j in range(length):
            ...
```

Trading Safety for performance

```
# cython: boundscheck = False
# cython: wraparound = False

def sum_arg(double[:] foo):
    ...
    ...

    with boundscheck(False), wraparound(False)
        for j in range(length):
            ...
```






**So what have we
learned?**

PROJECT DEMO
CYVLFEAT
GOOGLE SUMMER OF CODE 2016
@ PORTLAND STATE UNIVERSITY

A CYTHON / PYTHON WRAPPER FOR VLFEAT



WHAT IS VLFEAT?



WHAT IS CYVLFEAT?

[GITHUB.COM/MENPO/CYVLFEAT](https://github.com/menpo/cyvlfeat)

A man in a white soccer jersey and red shorts is celebrating with a young boy in a black soccer jersey and black shorts on a green soccer field. They are both smiling and running. In the background, another boy in a black hoodie and black shorts is walking. The scene is outdoors with a fence and trees in the background.

CONTRIBUTIONS ARE WELCOME!
[GITHUB.COM/SIMMIMOURYAL/CYVLFEAT](https://github.com/simmimouryal/cyvlfeat)

FEATURES WHICH NEED MORE WORK

1. BINSUM
2. KDTREE MODULE

github.com/simmimourya1/pycon_india_17

THANK YOU

